

```
In [1]: ## Name: Anmol Dhar
## Roll NO: 14113
## Subject:LP-IV(DL)
```

```
In [1]: import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, precision_score

RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]
```


```
In [2]: dataset = pd.read_csv("creditcard.csv")
```

```
In [3]: #check for any null values
print("Any nulls in the dataset",dataset.isnull().values.any())
print('-----')
print("No. of unique labels",len(dataset['Class'].unique()))
print("Label values",dataset.Class.unique())

#0 is for normal credit card transaction
#1 is for fraudulent credit card transaction
print('-----')
print("Break down of Normal and Fraud Transactions")
print(pd.value_counts(dataset['Class'],sort=True))
```

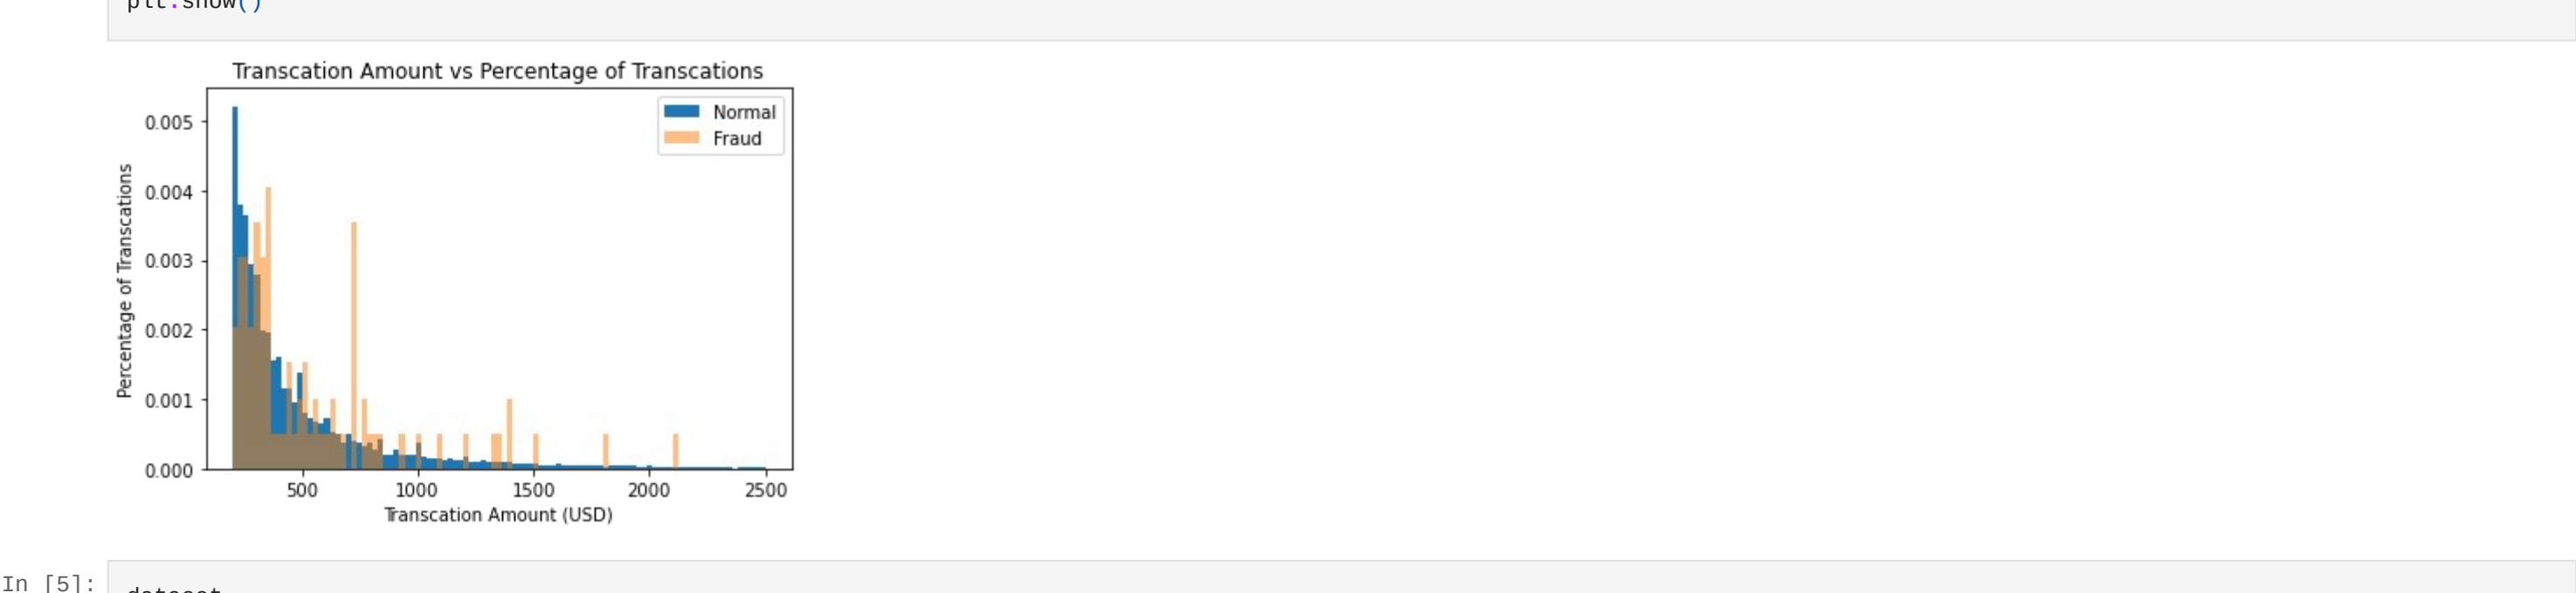
```
Any nulls in the dataset False
-----
No. of unique labels 2
Label values [0 1]
-----
Break down of Normal and Fraud Transactions
0      284315
1         492
Name: Class, dtype: int64
```

```
In [4]: #visualizing the imbalanced dataset
count_classes = pd.value_counts(dataset['Class'],sort=True)
count_classes.plot(kind='bar',rot=0)
plt.xticks(range(len(dataset['Class'].unique()))),dataset.Class.unique())
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
```



```
In [4]: #Save the normal and fraudulent transactions in separate dataframe
normal_dataset = dataset[dataset.Class == 0]
fraud_dataset = dataset[dataset.Class == 1]

#Visualize transaction amounts for normal and fraudulent transactions
bins = np.linspace(200,2500,100)
plt.hist(normal_dataset.Amount,bins=bins,alpha=1,density=True,label='Normal')
plt.hist(fraud_dataset.Amount,bins=bins,alpha=0.5,density=True,label='Fraud')
plt.legend(loc='upper right')
plt.title("Transaction Amount vs Percentage of Transactions")
plt.xlabel("Transaction Amount (USD)")
plt.ylabel("Percentage of Transactions")
plt.show()
```



In [5]: dataset

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.096698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	1.0	-1.358354	-1.340163	1.73209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	2.0	-1.58233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066556	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	0
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.640134	0.285745	-0.087371	0.004455	-0.026561	67.88	0
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	0
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	0

284807 rows x 31 columns

```
In [6]: sc = StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))
```

```
In [7]: raw_data = dataset.values
#The last element contains if the transaction is normal which is represented by 0 and if fraud then 1
labels = raw_data[:, -1]

#The other data points are the electrocardiogram data
data = raw_data[:,0:-1]

train_data, test_data, train_labels, test_labels = train_test_split(data, labels, test_size = 0.2, random_state =2021)
```

```
In [8]: min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data,tf.float32)
test_data = tf.cast(test_data,tf.float32)
```

```
In [9]: train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

#Creating normal and fraud datasets
normal_train_data = train_data[train_labels]
normal_test_data = test_data[~test_labels]

fraud_train_data = train_data[~train_labels]
fraud_test_data = test_data[test_labels]
print("No. of records in Fraud Train Data=",len(fraud_train_data))
print("No. of records in Normal Train Data=",len(normal_train_data))
print("No. of records in Fraud Test Data=",len(fraud_test_data))
print("No. of records in Normal Test Data=",len(normal_test_data))

No. of records in Fraud Train Data= 389
No. of records in Normal Train Data= 227456
No. of records in Fraud Test Data= 493
No. of records in Normal Test Data= 56859
```

```
In [10]: nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1]
#num of columns,30
encoding_dim = 14
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = 4
learning_rate = 1e-7
```

```
In [11]: #input layer
input_layer = tf.keras.layers.Input(shape=(input_dim,))

#Encoder
encoder = tf.keras.layers.Dense(encoding_dim,activation="tanh",activity_regularizer = tf.keras.regularizers.l2(learning_rate))(input_layer)
encoder = tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim2,activation=tf.nn.leaky_relu)(encoder)

#Decoder
decoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
decoder = tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim,activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim,activation='tanh')(decoder)

#Autoencoder
autoencoder = tf.keras.Model(inputs = input_layer,outputs = decoder)
autoencoder.summary()

Model: "model"
Layer (type) Output Shape Param #
-----
input_1 (InputLayer) [(None, 30)] 0
dense (Dense) (None, 14) 434
dropout (Dropout) (None, 14) 0
dense_1 (Dense) (None, 7) 105
dense_2 (Dense) (None, 4) 32
dense_3 (Dense) (None, 7) 35
dropout_1 (Dropout) (None, 7) 0
dense_4 (Dense) (None, 14) 112
dense_5 (Dense) (None, 30) 450
-----
Total params: 1,168
Trainable params: 1,168
Non-trainable params: 0
```

```
In [12]: cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5",mode='min',monitor='val_loss',verbose=2,save_best_only=True)
#Define our early stopping
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=1,
    mode='min',
    restore_best_weights=True
)
```

```
In [13]: autoencoder.compile(metrics=['accuracy'],loss= 'mean_squared_error',optimizer='adam')
```


```
In [14]: history = autoencoder.fit(normal_train_data,normal_train_data,epochs = nb_epoch,
                                batch_size = batch_size,shuffle = True,
                                validation_data = (test_data,test_data),
                                verbose=1,
                                callbacks = [cp,early_stop]).history
```

Epoch 1/50  
3543/3554 [=====>] - ETA: 0s - loss: 0.0033 - accuracy: 0.0372  
Epoch 1: val\_loss improved from inf to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 11s 2ms/step - loss: 0.0033 - accuracy: 0.0372 - val\_loss: 2.0179e-05 - val\_accuracy: 0.0343  
Epoch 2/50  
3545/3554 [=====>] - ETA: 0s - loss: 1.9583e-05 - accuracy: 0.0069  
Epoch 2: val\_loss did not improve from 0.00002  
3554/3554 [=====] - 6s 2ms/step - loss: 1.9599e-05 - accuracy: 0.0069 - val\_loss: 2.0190e-05 - val\_accuracy: 0.0078  
Epoch 3/50  
3517/3554 [=====>] - ETA: 0s - loss: 1.9562e-05 - accuracy: 0.0019  
Epoch 3: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 7s 2ms/step - loss: 1.9590e-05 - accuracy: 0.0619 - val\_loss: 2.0025e-05 - val\_accuracy: 0.0429  
Epoch 4/50  
3527/3554 [=====>] - ETA: 0s - loss: 1.9545e-05 - accuracy: 0.0599  
Epoch 4: val\_loss did not improve from 0.00002  
3554/3554 [=====] - 6s 2ms/step - loss: 1.9530e-05 - accuracy: 0.0601 - val\_loss: 2.0277e-05 - val\_accuracy: 0.2168  
Epoch 5/50  
3549/3554 [=====>] - ETA: 0s - loss: 1.8826e-05 - accuracy: 0.1758  
Epoch 5: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 6s 2ms/step - loss: 1.8831e-05 - accuracy: 0.1759 - val\_loss: 1.8344e-05 - val\_accuracy: 0.2184  
Epoch 6/50  
3539/3554 [=====>] - ETA: 0s - loss: 1.7526e-05 - accuracy: 0.2362  
Epoch 6: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 8s 2ms/step - loss: 1.7518e-05 - accuracy: 0.2363 - val\_loss: 1.7095e-05 - val\_accuracy: 0.3538  
Epoch 7/50  
3518/3554 [=====>] - ETA: 0s - loss: 1.8826e-05 - accuracy: 0.1125  
Epoch 7: val\_loss did not improve from 0.00002  
3554/3554 [=====] - 5s 2ms/step - loss: 1.8813e-05 - accuracy: 0.1137 - val\_loss: 1.7990e-05 - val\_accuracy: 0.2041  
Epoch 8/50  
3529/3554 [=====>] - ETA: 0s - loss: 1.7206e-05 - accuracy: 0.2141  
Epoch 8: val\_loss did not improve from 0.00002  
3554/3554 [=====] - 5s 1ms/step - loss: 1.7328e-05 - accuracy: 0.2143 - val\_loss: 1.7153e-05 - val\_accuracy: 0.2781  
Epoch 9/50  
3518/3554 [=====>] - ETA: 0s - loss: 1.6837e-05 - accuracy: 0.2481  
Epoch 9: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 6s 2ms/step - loss: 1.6825e-05 - accuracy: 0.2482 - val\_loss: 1.6811e-05 - val\_accuracy: 0.3518  
Epoch 10/50  
3516/3554 [=====>] - ETA: 0s - loss: 1.6692e-05 - accuracy: 0.2509  
Epoch 10: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
3554/3554 [=====] - 6s 2ms/step - loss: 1.6682e-05 - accuracy: 0.2510 - val\_loss: 1.6409e-05 - val\_accuracy: 0.3492  
Epoch 11/50  
3526/3554 [=====>] - ETA: 0s - loss: 1.6509e-05 - accuracy: 0.2484  
Epoch 11: val\_loss improved from 0.00002 to 0.00002, saving model to autoencoder\_fraud.h5  
Restoring model weights from the end of the best epoch: 1.  
3554/3554 [=====] - 7s 2ms/step - loss: 1.6561e-05 - accuracy: 0.2484 - val\_loss: 1.6237e-05 - val\_accuracy: 0.2865  
Epoch 11: early stopping

```
In [15]: plt.plot(history['loss'],linewidth = 2,label = 'Train')
plt.plot(history['val_loss'],linewidth = 2,label = 'Test')
plt.legend(loc='upper right')
plt.title('Model Loss')
plt.ylabel('Model Loss')
plt.xlabel('Epoch')

#plt.ylim(ymin=0.70,ymax=1)

plt.show()
```

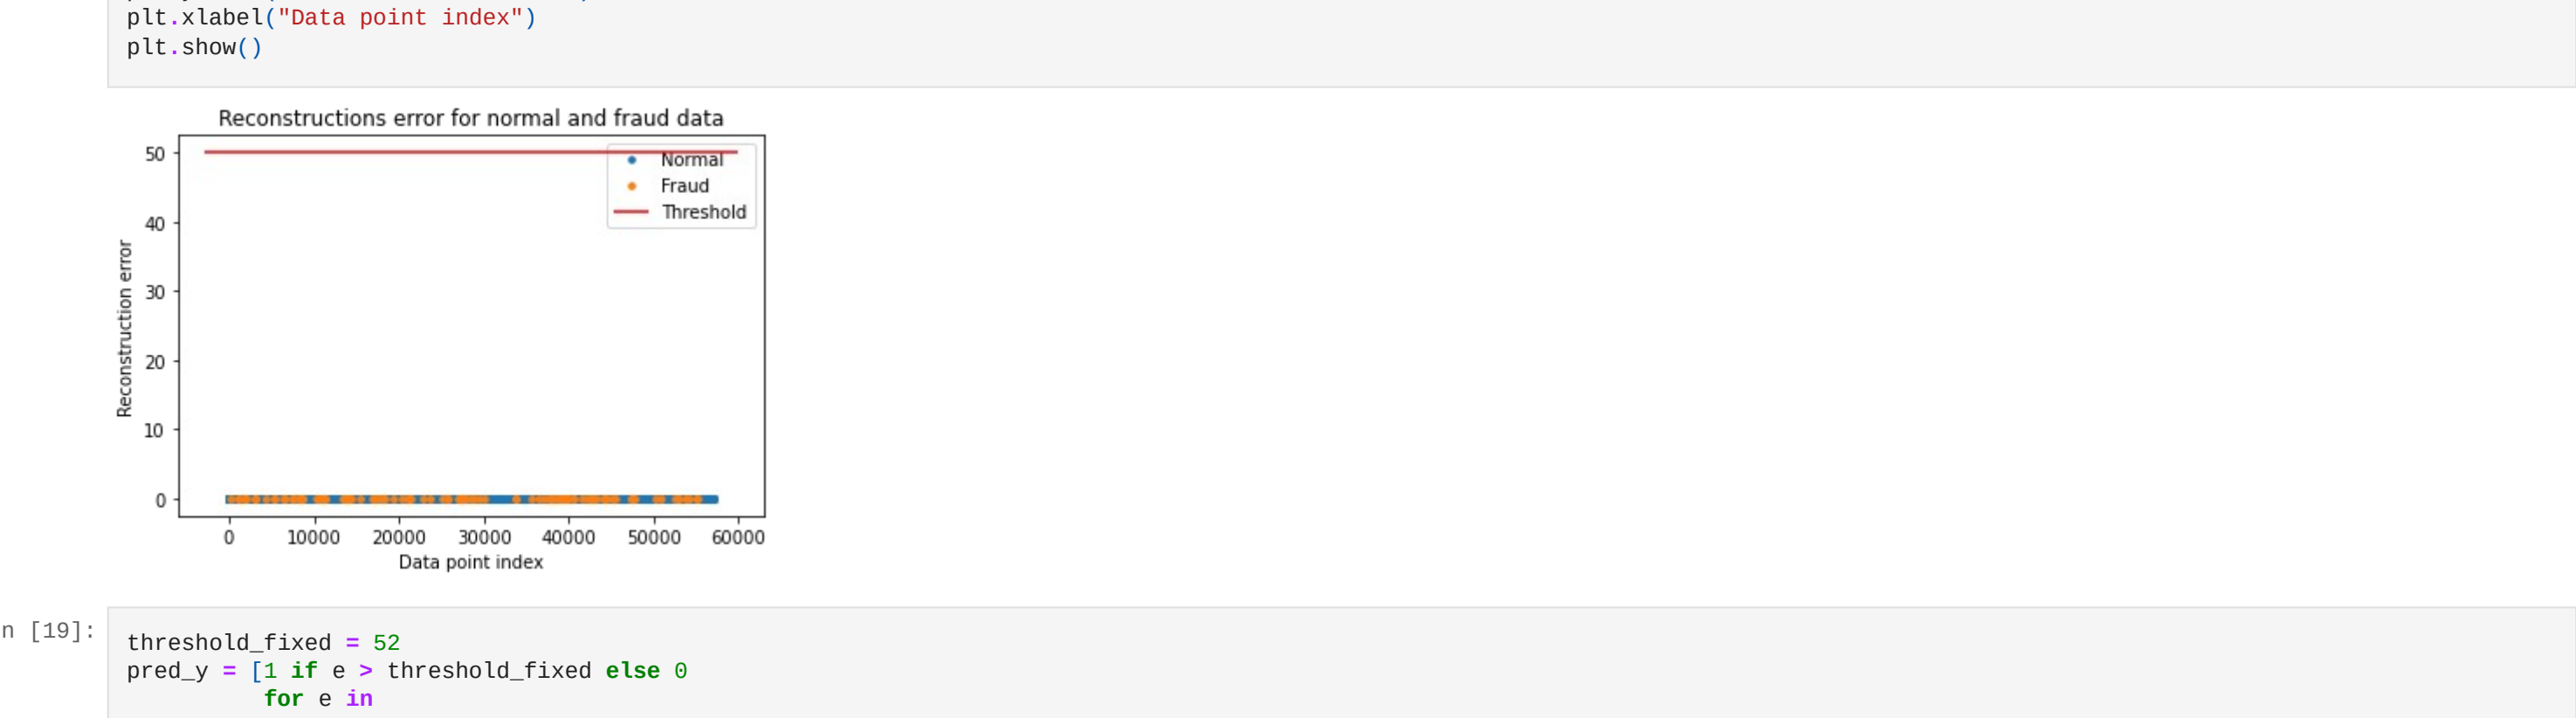


```
In [16]: test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2),axis = 1)
error_df = pd.DataFrame({'Reconstruction_error':mse,
                        'True_class':test_labels})
```

1781/1781 [=====] - 2s 865us/step

```
In [17]: threshold_fixed = 50
groups = error_df.groupby('True_class')
fig,ax = plt.subplots()


for name,group in groups:
    ax.plot(group.index,group.Reconstruction_error,marker='o',ms = 3.5,linestyle='',
            label = "Fraud" if name=="1" else "Normal")
ax.lines[threshold_fixed-1].get_xdata()[0],ax.get_xlim()[1],colors="r",zorder=100, label="Threshold")
ax.legend()
plt.title("Reconstructions error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show()
```



```
In [19]: threshold_fixed = 52
thresh_y = [1 if e > threshold_fixed else 0
            for e in
            error_df.Reconstruction_error.values]
error_df['pred'] = thresh_y
conf_matrix = confusion_matrix(error_df.True_class,pred_y)

plt.figure(figsize = (4,4))
sns.heatmap(conf_matrix,xticklabels = LABELS,yticklabels = LABELS,annot = True,fmt="d")
plt.title("Confusion matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()

#Print Accuracy,Precision and Recall
print("Accuracy :",accuracy_score(error_df['True_class'],error_df['pred']))
print("Recall :",recall_score(error_df['True_class'],error_df['pred']))
print("Precision :",precision_score(error_df['True_class'],error_df['pred']))
```



Accuracy : 0.9981917769741231  
Recall : 0.0  
Precision : 0.0